

Homework 6

PSTAT 131/231

Contents

Tree-Based Models	1
-----------------------------	---

Tree-Based Models

For this assignment, we will continue working with the file "pokemon.csv".

Each Pokémon has a primary type. Based on their type, a Pokémon is strong against some types, and vulnerable to others. (Think rock, paper, scissors.) A Fire-type Pokémon, for example, is vulnerable to Water-type Pokémon, but strong against Grass-type.

```
Pokemon <- read.csv("Pokemon.csv", header=TRUE)
head(Pokemon)
```

```
##   X.           Name Type.1 Type.2 Total HP Attack Defense Sp..Atk
## 1  1      Bulbasaur  Grass Poison   318 45    49    49    65
## 2  2      Ivysaur   Grass Poison   405 60    62    63    80
## 3  3      Venusaur  Grass Poison   525 80    82    83   100
## 4  3 VenusaurMega Venusaur  Grass Poison   625 80   100   123   122
## 5  4      Charmander   Fire         309 39    52    43    60
## 6  5      Charmeleon   Fire         405 58    64    58    80
##   Sp..Def Speed Generation Legendary
## 1     65    45           1      False
## 2     80    60           1      False
## 3    100    80           1      False
## 4    120    80           1      False
## 5     50    65           1      False
## 6     65    80           1      False
```

The goal of this assignment is to build a statistical learning model that can predict the **primary type** of a Pokémon based on its generation, legendary status, and six battle statistics.

Note: Fitting ensemble tree-based models can take a little while to run. Consider running your models outside of the .Rmd, storing the results, and loading them in your .Rmd to minimize time to knit.

Exercise 1

Read in the data and set things up as in Homework 5:

- Use `clean_names()`

- Filter out the rarer Pokémon types
- Convert `type_1` and `legendary` to factors

Do an initial split of the data; you can choose the percentage for splitting. Stratify on the outcome variable.

Fold the training set using v -fold cross-validation, with $v = 5$. Stratify on the outcome variable.

Set up a recipe to predict `type_1` with `legendary`, `generation`, `sp_atk`, `attack`, `speed`, `defense`, `hp`, and `sp_def`:

- Dummy-code `legendary` and `generation`;
- Center and scale all predictors.

Answer :

```
library(tidyverse)
library(tidymodels)
library(rpart.plot)
library(vip)
library(janitor)
```

```
Pokemon <-
  Pokemon %>%
  clean_names()
```

```
#str(Pokemon$type_1)
length(unique(Pokemon$type_1))
```

```
## [1] 18
```

```
filterdf <- Pokemon %>% filter(type_1%in% c("Bug", "Fire", "Grass", "Normal","Water","Psychic" )) %>%
  mutate(type_1 = factor(type_1),
         legendary = factor(legendary))
```

```
filterdf_split <- filterdf %>%
  initial_split(strata = type_1, prop = 0.7)
filterdf_train <- training(filterdf_split)
filterdf_test <- testing(filterdf_split)
dim(filterdf_train)
```

```
## [1] 318 13
```

```
dim(filterdf_test)
```

```
## [1] 140 13
```

```
Auto_folds <- vfold_cv(filterdf_train, strata = type_1, v = 5)
Auto_folds
```

```
## # 5-fold cross-validation using stratification
## # A tibble: 5 x 2
##   splits      id
##   <list>      <chr>
## 1 <split [252/66]> Fold1
## 2 <split [253/65]> Fold2
## 3 <split [253/65]> Fold3
## 4 <split [256/62]> Fold4
## 5 <split [258/60]> Fold5
```

```
filterdf_recipe <- recipe(type_1 ~ legendary+ generation + sp_atk +
                           attack + speed+ hp+sp_def, filterdf_train) %>%
  step_novel(all_nominal_predictors()) %>%
  step_dummy(all_nominal_predictors()) %>%
  step_zv(all_predictors()) %>%
  step_normalize(all_predictors())
```


Exercise 2

Create a correlation matrix of the training set, using the `corrplot` package. *Note: You can choose how to handle the continuous variables for this plot; justify your decision(s).*

What relationships, if any, do you notice? Do these relationships make sense to you?

```
library(corrplot)
library(corr)

filterdf_train %>%
  select(is.numeric, -x, -generation) %>%
  cor(use = "complete.obs") %>%
  corrplot(type = "lower", diag = FALSE)
```



special attack is positively correlated with damage resistance against special attacks and Speed is negatively correlated with defense and hit point

Exercise 3

First, set up a decision tree model and workflow. Tune the `cost_complexity` hyperparameter. Use the same levels we used in Lab 7 – that is, `range = c(-3, -1)`. Specify that the metric we want to optimize is `roc_auc`.

Print an `autoplot()` of the results. What do you observe? Does a single decision tree perform better with a smaller or larger complexity penalty?

Answer Decision tree performs better with small complexity penalty

```
tree_mod <-
  decision_tree(
    cost_complexity = tune(),
    tree_depth = tune()
  ) %>%
  set_engine("rpart") %>%
  set_mode("classification")
```

```
class_tree_wf <- workflow() %>%
  add_model(tree_mod) %>%
  add_recipe(filterdf_recipe)
```

```

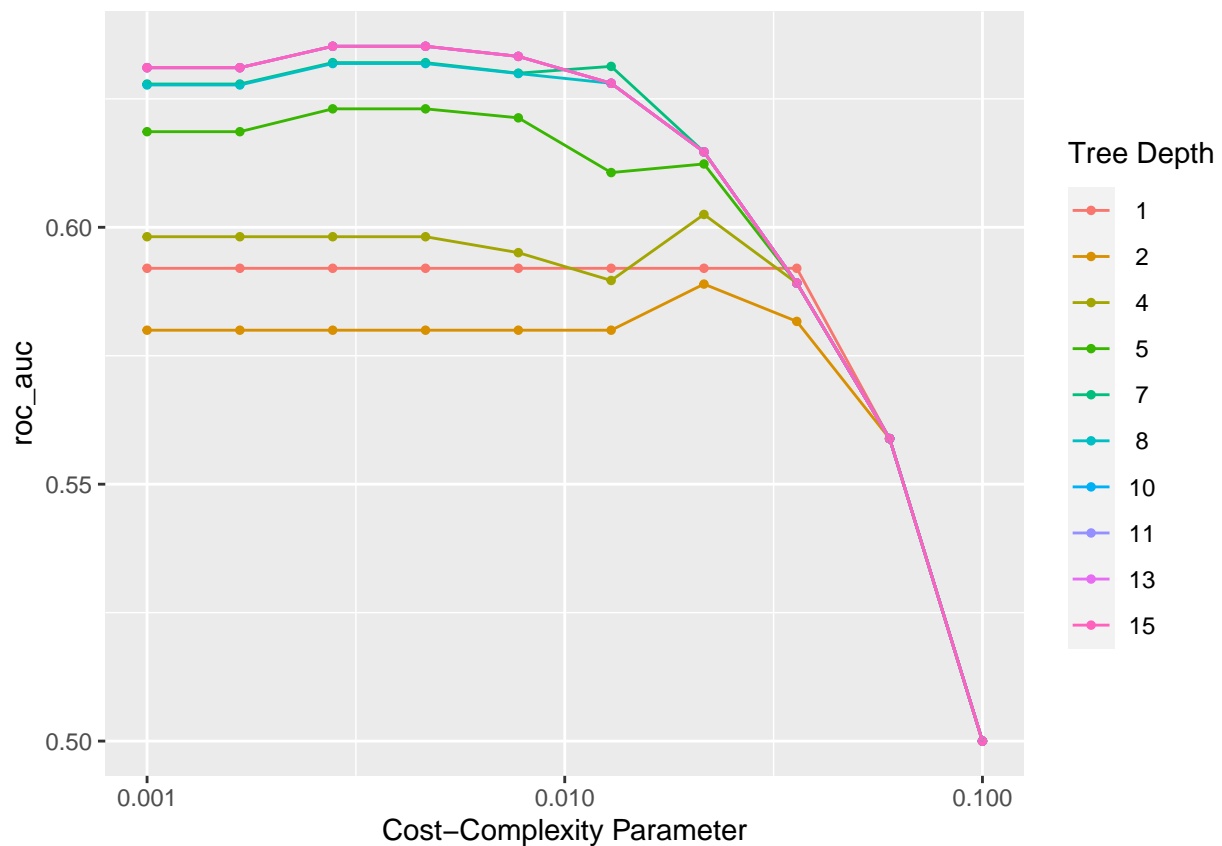
set.seed(3435)

param_grid <- grid_regular(cost_complexity(range = c(-3, -1)), tree_depth(), levels = 10)

tune_res <- tune_grid(
  class_tree_wf,
  resamples = Auto_folds,
  grid = param_grid,
  metrics = metric_set(roc_auc)
)

autoplot(tune_res)

```



Exercise 4

What is the `roc_auc` of your best-performing pruned decision tree on the folds? *Hint: Use `collect_metrics()` and `arrange()`.*

```

best_tree <- select_best(tune_res)
best_tree

## # A tibble: 1 x 3
##   cost_complexity tree_depth .config
##         <dbl>       <int> <chr>
## 1      0.00278         10 Preprocessor1_Model1063

```

```
collect_metrics(tune_res, summarize = TRUE) %>% arrange(select_best(tune_res), .by_group = TRUE)
```

```
## # A tibble: 100 x 8
##   cost_complexity tree_depth .metric .estimator mean      n std_err .config
##   <dbl>          <int> <chr>   <chr>    <dbl> <int>  <dbl> <chr>
## 1      0.001            1 roc_auc hand_till  0.592     5 0.00976 Preprocess~
## 2      0.00167          1 roc_auc hand_till  0.592     5 0.00976 Preprocess~
## 3      0.00278          1 roc_auc hand_till  0.592     5 0.00976 Preprocess~
## 4      0.00464          1 roc_auc hand_till  0.592     5 0.00976 Preprocess~
## 5      0.00774          1 roc_auc hand_till  0.592     5 0.00976 Preprocess~
## 6      0.0129          1 roc_auc hand_till  0.592     5 0.00976 Preprocess~
## 7      0.0215          1 roc_auc hand_till  0.592     5 0.00976 Preprocess~
## 8      0.0359          1 roc_auc hand_till  0.592     5 0.00976 Preprocess~
## 9      0.0599          1 roc_auc hand_till  0.559     5 0.0255 Preprocess~
## 10     0.1            1 roc_auc hand_till  0.5       5 0       Preprocess~
## # ... with 90 more rows
```

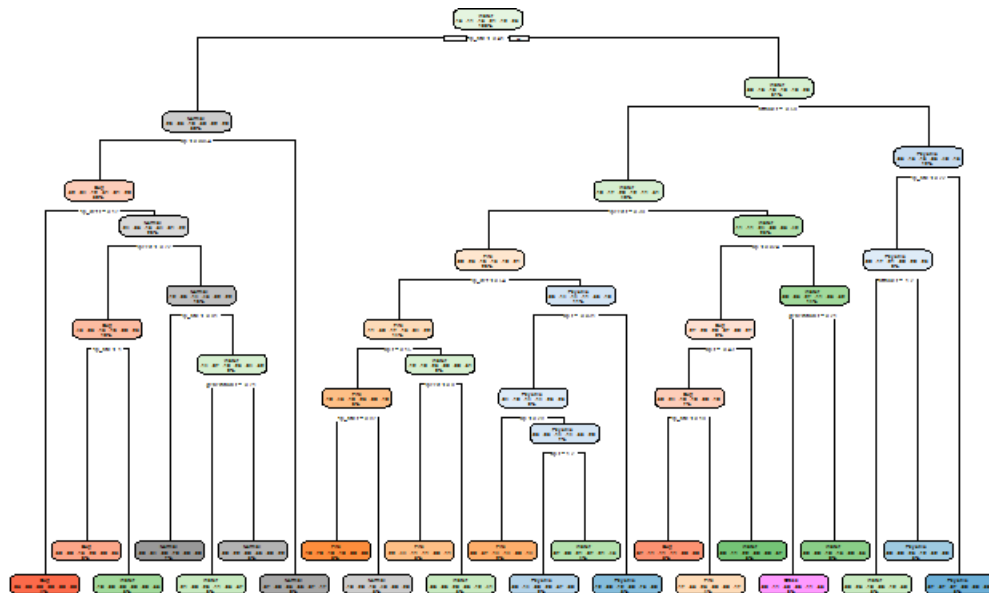
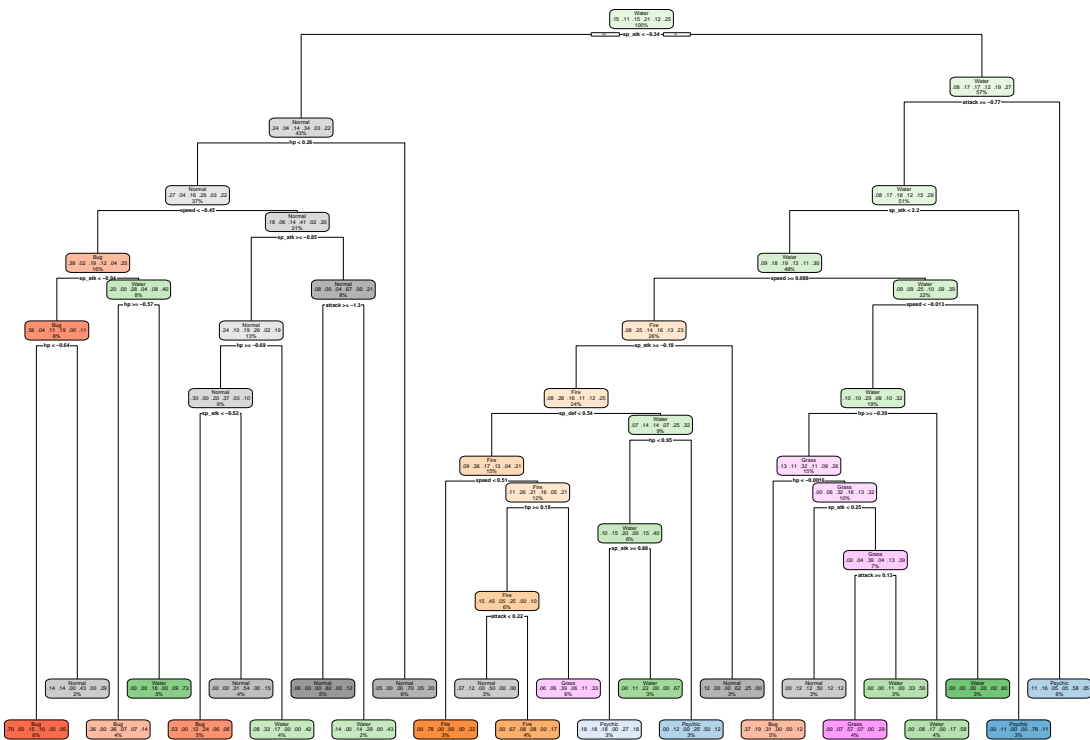
Exercise 5

Using `rpart.plot`, fit and visualize your best-performing pruned decision tree with the *training* set.

```
class_tree_final <- finalize_workflow(class_tree_wf, best_tree)

class_tree_final_fit <- fit(class_tree_final, data = filterdf_train)

class_tree_final_fit %>%
  extract_fit_engine() %>%
  rpart.plot()
```



Exercise 5

Now set up a random forest model and workflow. Use the `ranger` engine and set `importance = "impurity"`. Tune `mtry`, `trees`, and `min_n`. Using the documentation for `rand_forest()`, explain in your own words what each of these hyperparameters represent.

```
library(ranger)
cores <- parallel::detectCores()

rf_mod <-
  rand_forest(mtry = tune(), min_n = tune(), trees = 1000) %>%
  set_engine("ranger", importance = "impurity", num.threads = cores) %>%
  set_mode("classification")

random_tree_wf <- workflow() %>%
  add_model(rf_mod) %>%
  add_recipe(filterdf_recipe)
```

Create a regular grid with 8 levels each. You can choose plausible ranges for each hyperparameter. Note that `mtry` should not be smaller than 1 or larger than 8. **Explain why not. What type of model would `mtry = 8` represent?**

Exercise 6

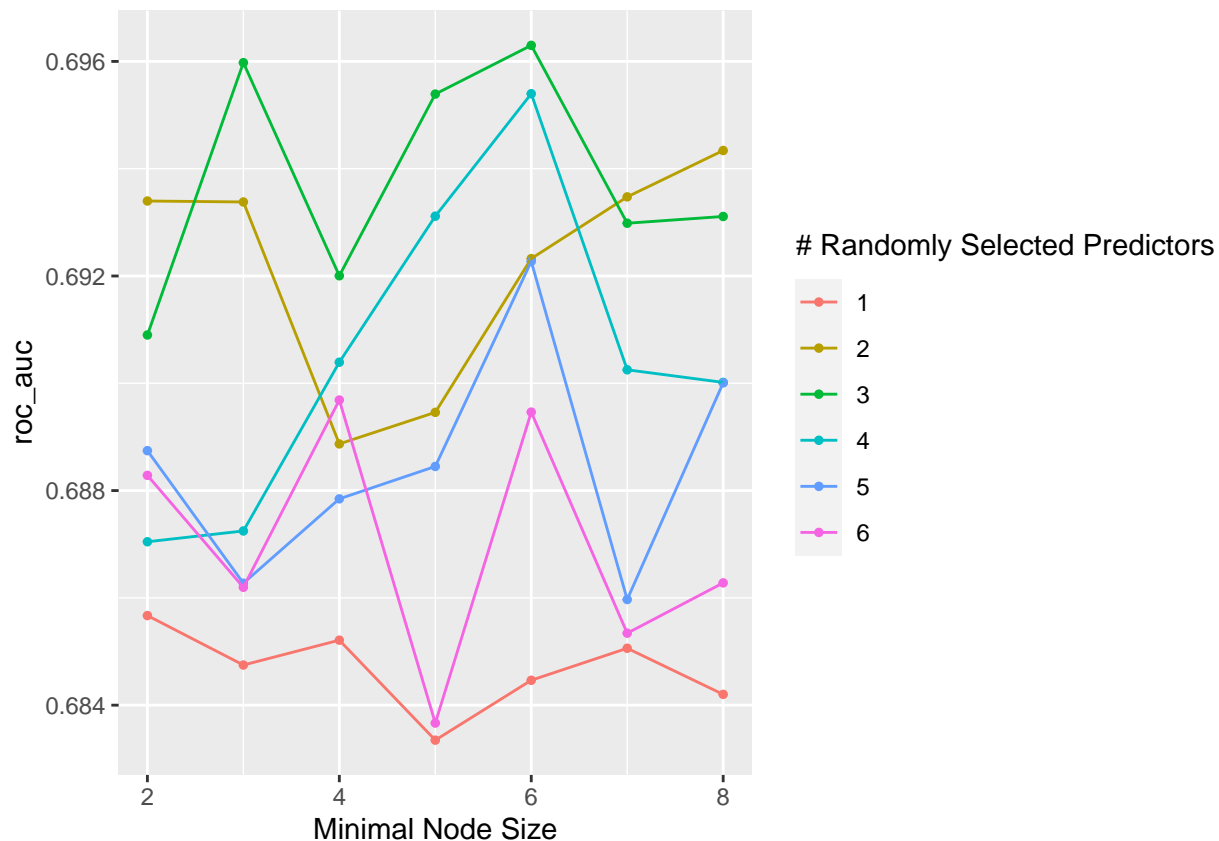
Specify `roc_auc` as a metric. Tune the model and print an `autoplot()` of the results. What do you observe? What values of the hyperparameters seem to yield the best performance?

```
Auto_folds <- vfold_cv(filterdf_train, strata = type_1, v = 5)

rf_grid <- grid_regular(
  mtry(range = c(1, 6)),
  min_n(range = c(2, 8)),
  levels = 8
)

set.seed(3456)
tune_res_rf <- tune_grid(
  random_tree_wf,
  resamples = Auto_folds,
  grid=rf_grid,
  metrics = metric_set(roc_auc)
)

autoplot(tune_res_rf)
```

Exercise 7

What is the roc_auc of your best-performing random forest model on the folds? *Hint: Use collect_metrics() and arrange().*

```
best_rf <- select_best(tune_res_rf)
best_rf
```

```
## # A tibble: 1 x 3
##   mtry min_n .config
##   <int> <int> <chr>
## 1     3     6 Preprocessor1_Model27
```

```
collect_metrics(tune_res_rf, summarize = TRUE) %>% arrange(desc(mean), .by_group = TRUE)
```

```
## # A tibble: 42 x 8
##   mtry min_n .metric .estimator mean    n std_err .config
##   <int> <int> <chr>    <chr>    <dbl> <int>   <dbl> <chr>
## 1     3     6 roc_auc hand_till 0.696     5 0.0120 Preprocessor1_Model27
## 2     3     3 roc_auc hand_till 0.696     5 0.0120 Preprocessor1_Model09
## 3     4     6 roc_auc hand_till 0.695     5 0.0107 Preprocessor1_Model28
## 4     3     5 roc_auc hand_till 0.695     5 0.0116 Preprocessor1_Model21
## 5     2     8 roc_auc hand_till 0.694     5 0.0140 Preprocessor1_Model38
## 6     2     7 roc_auc hand_till 0.693     5 0.0139 Preprocessor1_Model32
```

```
## 7      2      2 roc_auc hand_till 0.693      5 0.0112 Preprocessor1_Model102
## 8      2      3 roc_auc hand_till 0.693      5 0.0119 Preprocessor1_Model108
## 9      4      5 roc_auc hand_till 0.693      5 0.0128 Preprocessor1_Model122
## 10     3      8 roc_auc hand_till 0.693      5 0.0135 Preprocessor1_Model139
## # ... with 32 more rows
```

Exercise 8

Create a variable importance plot, using `vip()`, with your best-performing random forest model fit on the *training* set.

Which variables were most useful? Which were least useful? Are these results what you expected, or not?

```
rf_tree_final <- finalize_workflow(random_tree_wf, best_tree)

rf_tree_final_fit <- fit(rf_tree_final, data = filterdf_train)
```

```
library(vip)


best_auc <- select_best(tune_res_rf, "roc_auc")

final_rf <- finalize_model(
  rf_mod,
  best_auc
)
final_rf
```

```
## Random Forest Model Specification (classification)
##
## Main Arguments:
##   mtry = 3
##   trees = 1000
##   min_n = 6
##
## Engine-Specific Arguments:
##   importance = impurity
##   num.threads = cores
##
## Computational engine: ranger

rf_fit <- fit(final_rf, type_1 ~ ., data = filterdf_train)

vip(rf_fit)
```



Exercise 9

Finally, set up a boosted tree model and workflow. Use the `xgboost` engine. Tune `trees`. Create a regular grid with 10 levels; let `trees` range from 10 to 2000. Specify `roc_auc` and again print an `autoplot()` of the results.

What do you observe?

What is the `roc_auc` of your best-performing boosted tree model on the folds? *Hint: Use `collect_metrics()` and `arrange()`.*

```
library(xgboost)

xgb_mod <-
  boost_tree(
    trees = tune(),
    #min_n = tune(),
    mtry = tune()) %>%
  set_engine("xgboost") %>%
  set_mode("classification")

xgb_tree_wf <- workflow() %>%
  add_model(xgb_mod) %>%
  add_recipe(filterdf_recipe)
```

```

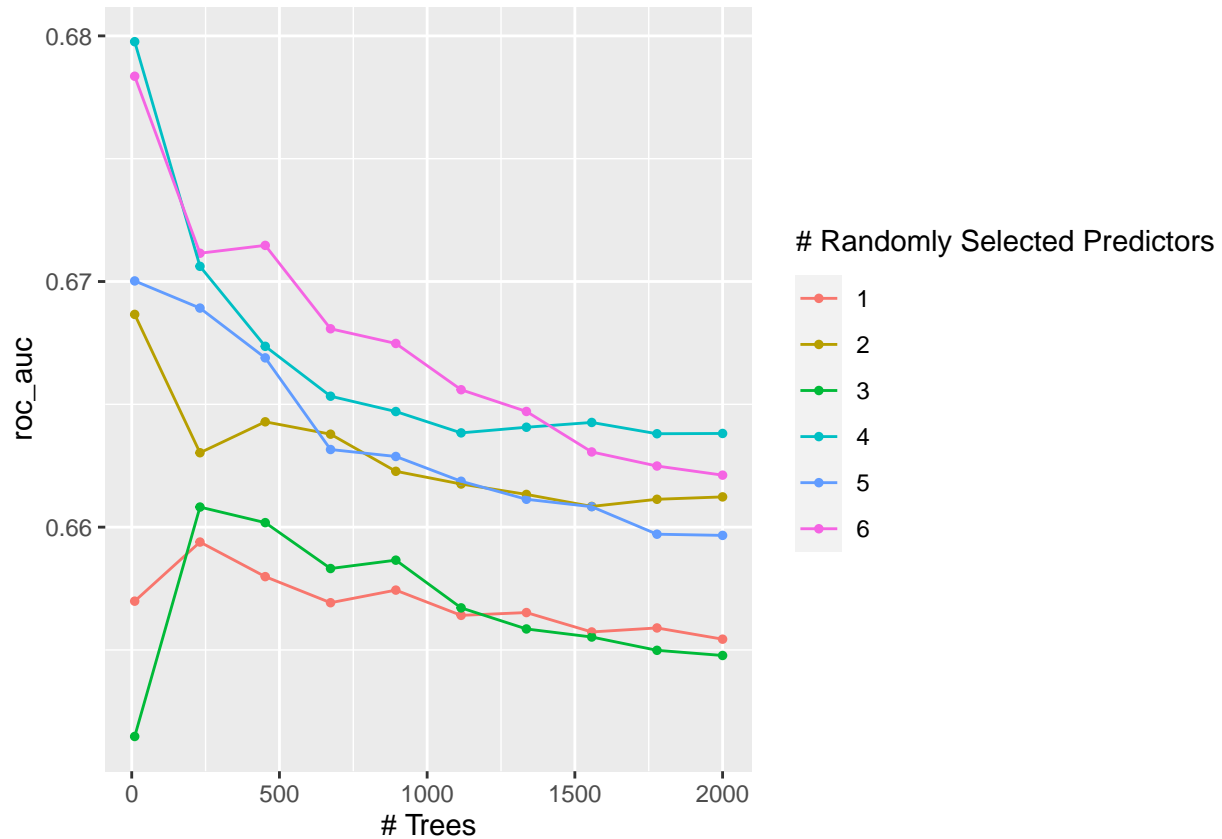
Auto_folds <- vfold_cv(filterdf_train, strata = type_1, v = 5)

xgb_grid <- grid_regular(
  mtry(range = c(1, 6)),
  #min_n(range = c(2,4)),
  trees(range = c(10,2000)),
  levels = 10
)

set.seed(3456)
tune_res_xgb <- tune_grid(
  xgb_tree_wf,
  resamples = Auto_folds,
  grid=xgb_grid,
  metrics = metric_set(roc_auc)
)

autoplot(tune_res_xgb)

```



```

best_xgb <- select_best(tune_res_xgb)
best_xgb

```

```

## # A tibble: 1 x 3
##   mtry trees .config
##   <int> <int> <chr>

```

```
## 1      4      10 Preprocessor1_Model31
```

```
collect_metrics(tune_res_xgb, summarize = TRUE) %>% arrange(desc(mean), .by_group = TRUE)
```

```
## # A tibble: 60 x 8
##   mtry trees .metric .estimator mean      n std_err .config
##   <int> <int> <chr>   <chr>    <dbl> <int>   <dbl> <chr>
## 1     4     10 roc_auc hand_till 0.680     5 0.0122 Preprocessor1_Model31
## 2     6     10 roc_auc hand_till 0.678     5 0.00879 Preprocessor1_Model51
## 3     6    452 roc_auc hand_till 0.671     5 0.0104 Preprocessor1_Model53
## 4     6    231 roc_auc hand_till 0.671     5 0.0111 Preprocessor1_Model52
## 5     4    231 roc_auc hand_till 0.671     5 0.0141 Preprocessor1_Model32
## 6     5     10 roc_auc hand_till 0.670     5 0.0147 Preprocessor1_Model41
## 7     5    231 roc_auc hand_till 0.669     5 0.0110 Preprocessor1_Model42
## 8     2     10 roc_auc hand_till 0.669     5 0.00839 Preprocessor1_Model11
## 9     6    673 roc_auc hand_till 0.668     5 0.00954 Preprocessor1_Model54
## 10    6    894 roc_auc hand_till 0.667     5 0.00929 Preprocessor1_Model55
## # ... with 50 more rows
```

```
xgb_tree_final <- finalize_workflow(xgb_tree_wf, best_xgb)
```

```
xgb_tree_final_fit <- fit(xgb_tree_final, data = filterdf_train)
```

Exercise 10

Display a table of the three ROC AUC values for your best-performing pruned tree, random forest, and boosted tree models. Which performed best on the folds? Select the best of the three and use `select_best()`, `finalize_workflow()`, and `fit()` to fit it to the *testing* set.

Print the AUC value of your best-performing model on the testing set. Print the ROC curves. Finally, create and visualize a confusion matrix heat map.

Which classes was your model most accurate at predicting? Which was it worst at?

```
tree_test_fit <- fit(class_tree_final, data = filterdf_train)
rf_tree_final_fit <- fit(rf_tree_final, data = filterdf_train)
xgb_tree_final_fit <- fit(xgb_tree_final, data = filterdf_train)
```

```
p1 <- augment(tree_test_fit, new_data = filterdf_train) %>%
  select(type_1, starts_with(".pred")) %>%
  roc_auc(type_1, .pred_Bug:.pred_Water)
```

```
p2<-augment(rf_tree_final_fit, new_data = filterdf_train) %>%
  select(type_1, starts_with(".pred"))%>%
  roc_auc(type_1, .pred_Bug:.pred_Water)
```

```
p3<-augment(xgb_tree_final_fit, new_data = filterdf_train) %>%
  select(type_1, starts_with(".pred")) %>%
  roc_auc(type_1, .pred_Bug:.pred_Water)
```

```
accuracies <- c(p1$.estimate,p2$.estimate,p3$.estimate)
models <- c("Decision Tree", "Random Forrest", "BoostedTree")
```

```
results <- tibble(roc = accuracies, models = models)
results %>%
  arrange(-accuracies)
```

```
## # A tibble: 3 x 2
##   roc models
##   <dbl> <chr>
## 1 0.998 BoostedTree
## 2 0.871 Decision Tree
## 3 0.5   Random Forrest
```

```
library(vip)

best_fit <- select_best(tune_res_xgb, tune_res_rf, tune_res, metric = "roc_auc" )

final_fit <- finalize_workflow(
  xgb_tree_final,
  best_fit
)

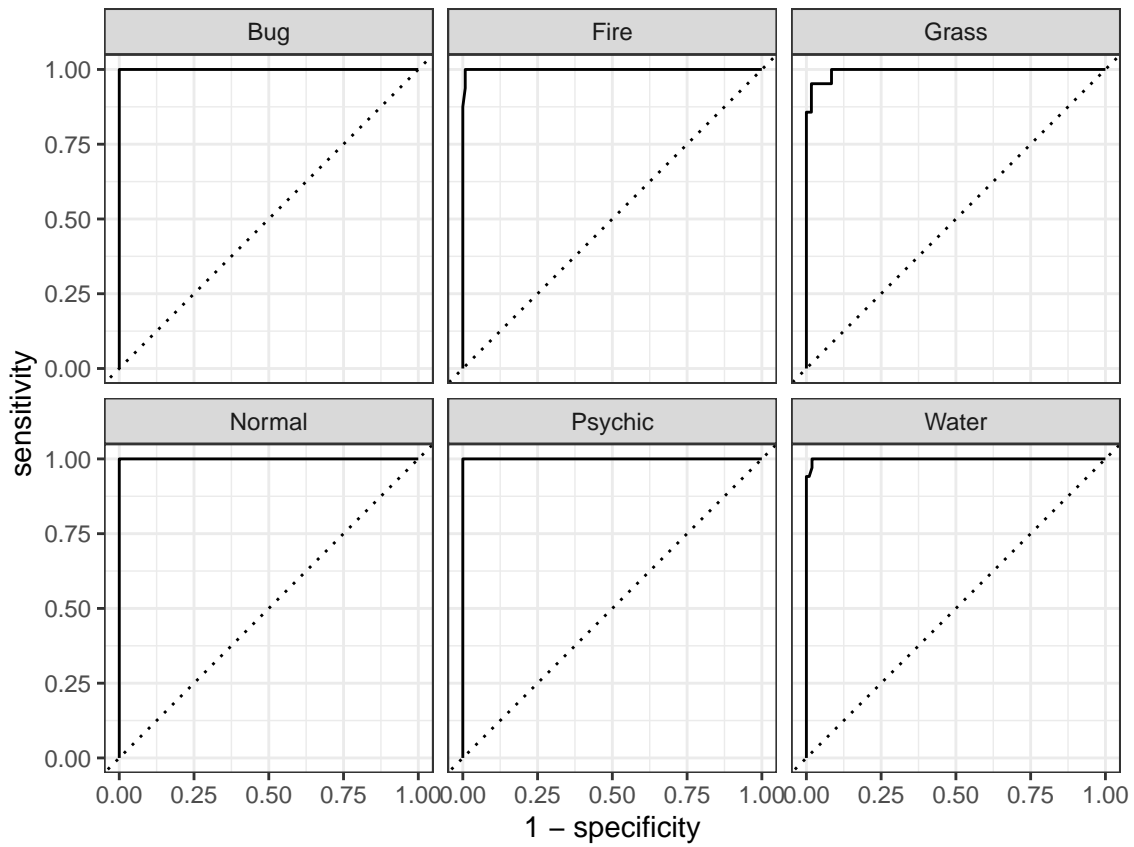
final_fit <- fit(final_fit, data = filterdf_test)

predicted_data <- augment(final_fit, new_data = filterdf_test) %>%
  select(type_1, starts_with(".pred"))

predicted_data %>% roc_auc(type_1, .pred_Bug:.pred_Water)
```

```
## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>         <dbl>
## 1 roc_auc hand_till     0.999
```

```
predicted_data %>% roc_curve(type_1, .pred_Bug:.pred_Water) %>%
  autoplot()
```



```
predicted_data%>%
  conf_mat(truth = type_1, estimate = .pred_class) %>%
  autoplot(type = "heatmap")
```

Prediction	Bug -	21	0	0	0	0	0
	Fire -	0	16	0	0	0	1
	Grass -	0	0	19	0	0	0
	Normal -	0	0	0	30	0	0
	Psychic -	0	0	1	0	18	0
	Water -	0	0	1	0	0	33
		Bug	Fire	Grass	Normal	Psychic	Water
		Truth					